

The Bloody Way to C

A Brutalist Approach to C Language—No Boilerplate

Simone Lungarella (simonelungarella@gmail.com)

2025-05-11

Contents

Preface	2
Introduction	3
1 Anatomy of a C program	4
1.1 Execute a C program	4
1.2 Include other source code	6
1.3 Functions	7
1.4 Variables	8
1.5 Code blocks	11
1.6 Iterating	13
1.7 Pointers	15

Preface

This e-book is a **humble** attempt to describe *C* language while actively trying to learn it. I enjoy writing code and technical documentation and I decided to produce this guide under [MIT licence](#). It is *not* intended to be fully comprehensive and complete, it only contains what I've learned and follows my very personal style.

It will be consistently updated and improved until completion and kept—as much as possible—accessible.

To understand every aspect of *C*, many tools will be used and all examples will refer to [CLI](#) commands. I will be using [Neovim](#) as text editor and operate on a Linux machine. The output of commands and all examples may differ from machine to machine but the concepts will hopefully remain valid.

I strongly believe that the best way to develop software is by using *CLI* and lightweight text editors such as `neovim` or `vim`. Whenever it is possible, I will avoid using browsers to search for documentation by preferring usage of `man` directly into the terminal. This will keep low the friction and avoid the necessity to leave the home row of my keyboard.

Introduction

C was invented in [Bell Labs](#) when [Ken Thompson](#) was working on Unix. Following the idea that a good operating system should have had a high level compiled language. After abandoning the first attempt on creating a compiler for [Fortran](#), a smaller new language was created and named [B](#). [B](#) better fitted [P2P11](#) but was not enough to port Unix from Assembly. [C](#) was created with a set of feature that were missing in [B](#) and was a much better fit for the Unix system.

[C](#) was a better language mainly because its multiple distinct types:

- pointers;
- integer;
- floating point numbers: float;

In that sense, [C](#) language can be visualized as [B](#) with types where all types can also be imagined as integers since pointers—in very simple terms—are integers and so are structures. In fact, structures are a set of integers representing offsets of each field position in memory and values of the very same fields. This simplicity can be considered the strength of the language as it can be easily picked up by new developers, layered to build a powerful abstraction and, with that, imagine in simple terms complex topics and algorithms.

1 Anatomy of a C program

1.1 Execute a C program

C is a compiled language, this means that you cannot execute a file containing the main function. It requires to be compiled in an executable program.

You can use: `cc` to compile a *C* program. `cc` is a Unix command that let you easily communicate with the compiler. You can use: `cc --version` to check what compiler does it use.

```
cc (GCC) 14.2.1 20240912 (Red Hat 14.2.1-3)
Copyright (C) 2024 Free Software Foundation, Inc.
This is free software; see the source for copying conditions.  There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

Let's consider the following simple *C* program contained in a file named—for instance—*hello_world.c* (how original):

```
#include <stdio.h>

int main(void) {
    printf("Hello World\n");
    return 0;
}
```

You can use: `cc hello_world.c` to compile it to an executable program.

The compiler generates an executable binary file named: *a.out*. This file is executable and runs your program. You can use: `file a.out` to check information about the generated file.

```
./a.out: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked, \
interpreter /lib64/ld-linux-x86-64.so.2, \
BuildID[sha1]=d589730d718a032a35f848fe8d280063a6cee18c, \
for GNU/Linux 3.2.0, not stripped
```

If you want to check the content of the generated binary file, you can use: `hexdump -C a.out`.

You can also generate [Assembly](#) code using: `cc -S hello_world.c` if the compiler supports this feature.

```
.file    "hello_world.c"
.text
.section    .rodata
.LC0:
.string    "Hello World"
.text
.globl    main
.type     main, @function
main:
.LFB0:
.cfi_startproc
pushq    %rbp
.cfi_def_cfa_offset 16
.cfi_offset 6, -16
movq     %rsp, %rbp
.cfi_def_cfa_register 6
movl     $.LC0, %edi
call     puts
movl     $0, %eax
popq     %rbp
.cfi_def_cfa 7, 8
ret
.cfi_endproc
.LFE0:
.size    main, .-main
.ident   "GCC: (GNU) 14.2.1 20240912 (Red Hat 14.2.1-3)"
.section    .note.GNU-stack,"",@progbits
```

With a given compiler, you can tweak many compilation aspects. For instance, `cc -O2 hello_world.c` tells the compiler to optimize the generated executable. A more optimized version of the executable is also slower to be generated and, if that does not make much sense for small programs, it can output a much better version of the program when a high enough level of complexity has been reached.

With [GCC](#) compiler, you can see that our simple program make use of `puts` [syscall](#), however, this depends on the compiler itself and, often, with different compilers, the line: `printf("Hello World\n");` is compiled using `printf` instead.

Using `-O2` flag can make the compiler use `puts` as this syscall is faster than `printf`. This is a simple, yet meaningful, example but in such a small program it does make no difference in terms of execution speed. The compiler is very good at improving written programs if given enough time. While developing, though, a low compilation time is often preferred.

You can check the standard C library from the terminal using `man` or `--help` flags. For example, you can use `man 3 puts` or `man 3 printf` to check documentation of both syscalls (3 makes sure to output the C library description).

1.2 Include other source code

In the very first line of our simple program, you can see a [preprocessor directive](#). This line simply tells to the compiler that a file need to be included into the program. The compiler, before the compilation, take the content of the file and *paste* it at the location. In this case, `<stdio.h>` declares the prototype of `printf` function so to instruct the compiler on how to execute that specific call. To prove this point, you can remove the first line and replace it with: `int printf(const char *restrict format, ...);` which is the prototype of the function we want to call.

```
int printf(const char *restrict format, ...);

int main(void) {
    printf("Hello World\n");
    return 0;
}
```

`#include` can also be used to include other *C* files. In fact, you can move a single line to a different file and than compile a program that includes the file on the line you want it to be replaced.

```
#include <stdio.h>

int main(void) {
    #include "file.c"
    return 0;
}
```

The generated assembly or machine code will be equivalent.

1.3 Functions

This very simple program has a single function named *main*. A function has always a return type, an *optional* list of parameters, and a body. The signature of the function *main* has a return type specified as `int`—this means that the function must return an integer value.

Parameters are defined inside the brackets of the function and they too have a specific type. It is also possible to define a function that does not require any parameter. This can be explicit, using `void` as the function `main` does, or implicit, by simply avoiding specifying any parameter: `int main() {}`.

Functions can call other functions too!

```
#include <stdio.h>

int sum(int a, int b) {
    return a + b;
}

int main(void) {
    printf("Hello World %d\n", sum(10, 20));
    return 0;
}
```

The function `main` is a special function, in fact, it is the only function that is automatically called by the program. Other functions must be explicitly called. This means that a valid *C* program must define the *main* function.

1.4 Variables

Functions parameters are variables existing only during the function execution. There are variables which are not involved only in function calls but also have a meaning in the callee context or even in the whole program context.

1.4.1 Scope

Variables can have different scope. In the previous example, the function `int sum(int a, int b)` has two variables as parameters having a local scope. When variables are local, they are valid only within the function context and have no meaning to other functions.

To understand this concept, let's consider the following program:

```
#include <stdio.h>

int sum(int a, int b) {
    return a + b;
}

int main(void) {
    int a = 10;
    int c = 20;

    printf("Hello World %d\n", sum(a, c));
    return 0;
}
```

This is a valid *C* program, equivalent to the previous, and, as you can see, the variable named `a` exists twice with the same name. This is possible because in both cases, the variable scope is local to the function itself and it's removed after the function has returned its value.

The function `main` has a return type but since it is automatically called by the program, the only one that can be interested in its value is the callee: the program executor. If executed from a shell, the program returns its value and can be shown with `./a.out; echo $?`. This is quite useful combined with the fact that `0` is equivalent to `true` in Unix shells.

Variables can also have a global scope. A global variable is seen by every function and initialized only once.

```
#include <stdio.h>

int x = 0;

void incr(void) {
    x = x + 1;
    printf("%d\n", x);
}

int main(void) {
    incr();
    incr();

    return 0;
}
```

In such cases the value of `x` is incremented by one each time the function is called. Values of local variables can also be retained through multiple function calls if they are defined as static: `static int x = 0;`.

It's important to highlight that, in *C*, variables are passed **by value**. This means that whenever a function is called, it cannot modify any existing variable local to the callee but, for each of its parameters, a copy of the value is passed. To modify local variables with functions it is necessary to use *pointers* which will be described extensively in a dedicated chapter.

1.4.2 Type

We have seen variables having type *int*, but there are multiple primitive types that can define different kinds of data. For simplicity, a subset of common primitive types is shown into the following table, refer to the [standard documentation](#) to explore all different existing types.

Type	Common size (b)	Description
int	32	Signed integer numbers
float	32	Floating point numbers
double	64	High precision floating point numbers
char	8	Characters
short	16	Shorter signed integer numbers

All size reported are not guaranteed by C specification, it mainly depends on architectures.

In many cases, types are automatically *promoted* to a higher size type to easily handle similar cases. For instance, `printf` will promote `char` or `short` values to `int` enabling developers to simplify the usage of the function.

```
short s = 400;

// `s` is automatically converted.
printf("%d\n", s);
```

This happens with functions such as `printf`, which accept a variable number of parameters (variadic function), but also during expressions evaluations if necessary.

```
char c = 127;

// Before evaluation, `c` is promoted to int.
int i = c + 1;
printf("%d\n", i);
```

Since the size of types is variable and depends on the architecture, there is a specific function that returns the size of a specific variable: `sizeof(var)`.

Variables can represent a single value or a collection of values. To define a variable and store multiple values of the same type, *C* provides **Arrays**.

```
int array[5] = {1, 2, 3, 4, 5};

printf("%d\n", array[0]);
```

Arrays can store multiple values in different positions and track a contiguous block of memory. To access value in a specific position you can use each index starting from `0` up to `n-1`.

Arrays of characters are named *strings*, and—since they are very common—there is a simpler way to initialize them:

```
char phrase[] = "Hello World";

printf("%s\n", phrase);
```

It is not mandatory to set the size of the array as the compiler will automatically do it by checking the its size. You can evaluate the size of a string too using: `sizeof(string)`, which is also evaluated at compile time.

You will notice that the size of strings, returned by `sizeof`, is always greater than the amount of character in the string. Strings always require the null terminator: `\0` that tells the program when the array ends and initializing a string using quotes automatically adds the null terminator.

1.5 Code blocks

Code blocks are blocks delimited by brackets that can be part of functions. Each function has at least one code block. Variables declared in a specific block have a local meaning and occupy a different memory block.

```
#include <stdio.h>

int main(void) {
    int i = 5;

    {
        int i = 3;

        // (4 bytes) stored at 0x7fff7cae44b8
        printf("(%zu bytes) stored at %p\n", sizeof(i), &i);
    }

    // (4 bytes) stored at 0x7fff7cae44bc
    printf("(%zu bytes) stored at %p\n", sizeof(i), &i);
}
```

`&i` returns the memory address where the variable is stored, more about *pointers* in the following chapters.

This simple program will show how the two variables, having the same name, will be stored in two consecutive memory blocks that differ by exactly 4 bytes (from 0x[...]8 to 0x[...]c).

1.5.1 Conditional code blocks

Often, the linear execution of a program needs to be interrupted to take a direction based on a specific condition. Conditional code blocks are blocks of code executed only if a specific requirement is satisfied.

The keyword `if` defines a conditional block and the condition that needs to be met for the execution:

```
#include <stdio.h>

int main(void) {
    int i = 5;

    if (i > 3) {
        printf("Value %d is greater than 3\n", i);
    } else {
        printf("Value %d is not greater than 3\n", i);
    }
}
```

Conditional blocks are optionally enhanced with `else` or multiple `else if` constructs that build up the logic based on multiple different conditions.

When conditional code blocks are constituted by a single statement, brackets are optional:
`if (i > 0) printf(i); .`

1.5.2 Switch

Another useful way to handle conditional blocks is by using `switch` keyword.

```
#include <stdio.h>

int main(void) {
    int i = 5;

    switch(i) {
    case 5:
        printf("It is a five!");
        break;
    case 3:
        printf("It's not a five.");
        break;
    default:
        printf("No case matched.");
        break;
    }
}
```

Switch blocks can be used to execute code based on a matching condition in an elegant way. Each case code block have to be terminated with the `break` keyword to avoid executing following code blocks too.

1.6 Iterating

To be [Turing-complete](#), a language must have some kind of looping logic. *C* has many way to iterate the execution of a code block: `for`, `while`, `do-while`. Loops let the program jump at the start of a code block for multiple iteration each time the condition is met. A way to achieve the same result is by using the keyword `goto`.

The keyword `goto` interrupt the program execution and start from a specified *label*.

```
#include <stdio.h>

int main(void) {
    int i = 0;

again:
    printf("%d", i);
    i++;
    if (i < 10) goto again;
}
```

In the example, when the condition is met, the instruction `goto` make the program jump back to the line under the specified label.

The logic to iterate a code block or a set of instruction can be also written by using a *while loop*.

```
#include <stdio.h>

int main(void) {
    int i = 0;

    while(i < 10) {
        printf("%d", i);
        i++;
    }
}
```

1.6.1 Recursion

Another way to execute multiple times a specific code block, is by using recursion. We talk about recursion when a function call itself. In the following example, `count` is a recursive function.

```
#include <stdio.h>

void count(int start, int end) {
    if (start > end) return;
    printf("%d\n", start);
    count(start + 1, end);
}

int main(void) {
    count(0, 9);
}
```

```
    return 0;
}
```

Calling multiple time the same functions is equivalent to have multiple code blocks and—as said—each code block instantiate its variables in different memory addresses. This means that recursion, by instantiating multiple time the same variables, uses more memory than a simple while loop.

The recursive function *count* can be written using `goto` instead of recursion.

```
void count(int start, int end) {
    if (start > end) return;
iterate:
    printf("%d\n", start);
    start++;
    goto iterate;
}
```

1.6.2 For loops

Another way to iterate over elements is by using a *for loop*.

```
#include <stdio.h>

int main(void) {
    for(int i = 0; i < 10; i++) {
        printf("%d", i);
    }
}
```

This kind of loop defines the starting value of the loop counter, the exit condition and the increment rule on the same line, then defines the code to be executed in a code block. Loop counter, exit condition and increment rule can be omitted and handled manually if necessary.

```
#include <stdio.h>

int main(void) {
    int i = 0;

    for(;;) {
        if (i < 10) break;
        printf("%d", i);
        i++;
    }
}
```

It's important to notice that if the loop counter is defined the same time it is initialized, the variable will be local to the loop block.

1.7 Pointers

Pointers are special variables that indicate an area of the memory of a specific type. They are declared with an asterisk as in the expression: `int *y;` where `int` denotes the type of the data allocated at the address. Pointers can map an address of any type and have always the size of a single [word](#).

Every variable has its own address in memory, to get the address of a given variable, can be used the operator: `&`.

```
#include <stdio.h>

int main(void) {
    int x = 5;
    int *y = &x;

    printf("x is stored at the address: %p\n", y);
    return 0;
}
```

Values of variables can be modified interacting with its pointer too, so `int *y = 10` would alter the value of the variable `x`.

Pointers are extremely powerful because they make possible updating variables without any extra memory allocation. It is possible, in fact, to update a variable by calling a function that does not instantiate any local variable but only access and manipulate the value using its pointer.

```
#include <stdio.h>

void incr(int *p) {
    *p = *p + 1;
}

int main(void) {
    int x = 5;
    int *y = &x;

    printf("x was: %d\n", x);

    // This instruction alters the value of x using its pointer
    incr(y);
    printf("x now is: %d\n", x);
    return 0;
}
```

Pointers have always the size of a `intptr_t` as it represents a memory address. The type declared with the pointer indicates the type of the data stored at the given address.